

# Module 2-3

Joins

# Objectives

- Normalization
- Keys (Primary, Natural, Surrogate, Foreign)
- Cardinality (1-1, 1-M, M-M)
- SQL Joins (Inner and Left Join)
- Unions
- Create a new database (MovieDB)

A decorative graphic consisting of a thin gold circle on the left side. A thick black left square bracket is positioned to the left of the circle's center. A thick gold right square bracket is positioned to the right of the circle's center. A horizontal gold bar with a gradient from dark to light extends from the left bracket across the top of the slide.

# Normalization

# Normalization

- Normalization is a process by which data structures in a relational database are as efficient as possible, including the elimination of redundancy, the minimisation of the use of null values and the prevention of the loss of information.

# Aims of Normalization

- Normalization ensures that the database is structured in the best possible way.
- To achieve control over data redundancy. There should be no unnecessary duplication of data in different tables.
- To ensure data consistency. Where duplication is necessary the data is the same.
- To ensure tables have a flexible structure. E.g. number of classes taken or books borrowed should not be limited.
- To allow data in different tables can be used in complex queries.

# Duplication vs Redundant Data

Duplicated Data: When an attribute

- has two or more identical values

Redundant Data: If you can delete

- data with a loss of information

# Stages of Normalization

- First Normal Form (1NF)
- Second Normal Form (2NF)
- Third Normal Form (3NF)

## Candidate Key / Primary Key

- A *Candidate Key* is an attribute that can be used to uniquely identify each tuple (a single row in a table) in a relation.
- A relation may have more than one candidate key
- If so, one candidate key is nominated as the primary key




# [ First Normal Form ]

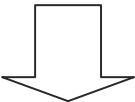
**A table is in its first normal form if  
it contains no repeating  
attributes or groups of  
attributes**

# Non-Normalised Table

STUDENT



 <u>Number</u>	Name	Classes
001231	William Hartnell	Information Systems, Systems Analysis, Data Communications
001232	Patrick Troughton	Systems Analysis, Data Communications
001233	Jon Pertwee	OO Programming, Systems Analysis, Data Communications
001234	Tom Baker	Systems Analysis, Data Communications

# First Normal Form

- To convert data for unnormalised form to 1NF, simply convert any repeated attributes into part of the candidate key
  - STUDENT(Number, Name, Classes)
- 
- STUDENT(Number, Name, Classes)

# First Normal Form

STUDENT

 <u>Number</u>	Name	 <u>Classes</u>
001231	William Hartnell	Information Systems
001231	William Hartnell	Systems Analysis
001231	William Hartnell	Data Communications
001232	Patrick Troughton	Systems Analysis
001232	Patrick Troughton	Data Communications
001233	Jon Pertwee	OO Programming
001233	Jon Pertwee	Systems Analysis
001233	Jon Pertwee	Data Communications
001234	Tom Baker	Systems Analysis
001234	Tom Baker	Data Communications

[ Over to you... ]

RefNo	Name	Address	Status	AccNo
345	C.J. Date	23, High Street	Business	120768, 348973
543	F.D. Rolland	45, The Ash	Domestic	987654
675	D.R. Howe	17, Low Street	Business	745363, 678453, 348973

# **Second Normal Form**

**A table is in the second normal form if  
it's in the first normal form AND no  
column that is not part of the primary  
key is dependant only a portion of the  
primary key**

# [ Second Normal Form ]

---



- The concept of functional dependency is central to Normalization and, in particular, strongly related to 2NF.

# Functional Dependency



- If 'X' is a set of attributes within a relation, then we say 'A' (an attribute or set of attributes), is functionally dependant on X, if and only if, for every combination of X, there is only one corresponding value of A
- We write this as :  
$$X \rightarrow A$$




# Table in 1NF

 RefNo	Name	Address	Status	 AccNo
345	C.J. Date	23, High Street	Business	120768
345	C.J. Date	23, High Street	Business	348973
543	F.D. Rolland	45, The Ash	Domestic	987654
675	D.R. Howe	17, Low Street	Business	745363
675	D.R. Howe	17, Low Street	Business	678453
675	D.R. Howe	17, Low Street	Business	348973

# Second Normal Form


 RefNo	 AccNo
345	120768
345	348973
543	987654
675	745363
675	678453
675	348973

 RefNo	Name	Address	Status
345	C.J. Date	23, High Street	Business
543	F.D. Rolland	45, The Ash	Domestic
675	D.R. Howe	17, Low Street	Business



[ Over to you... ]

 Supplier#	 Part#	City	Quantity
S1	P1	London	1000
S1	P2	London	1500
S1	P3	London	3400
S1	P4	London	2100
S2	P2	Paris	3400
S2	P3	Paris	1000
S4	P1	Nuku alofa	5
S4	P4	Nuku alofa	7

# Table in Second Normal Form



Supplier#	City
S1	London
S2	Paris
S4	Nuku alofa



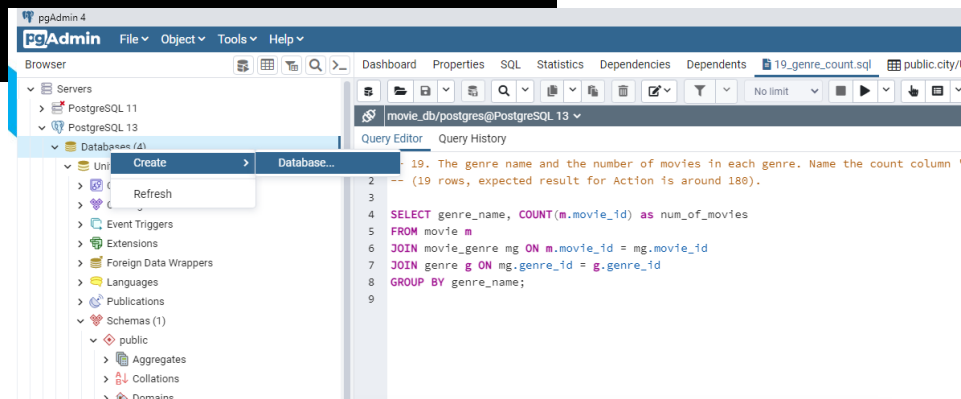
Supplier#	Part#	Quantity
S1	P1	1000
S1	P2	1500
S1	P3	3400
S1	P4	2100
S2	P2	3400
S2	P3	1000
S4	P1	5
S4	P4	7

## Third Normal Form

**A table is in the third normal form if it is the second normal form and there are no *non-key columns* dependant on *other non-key columns* that could not act as the primary key.**


# Create the MovieDB and insert some data

```
/usr/bin/bash --login -i
~/workspace/NLR-3/te-curriculum/module-2/03_Joins/lecture-student/postgres [release_v2_4]
17:08 $ winpty createdb -U postgres MovieDB
Password:
~/workspace/NLR-3/te-curriculum/module-2/03_Joins/lecture-student/postgres [release_v2_4]
+ 2]
17:08 $ psql -U postgres -d MovieDB -f MovieDB-data.psql |
```



# Amazon Scenario

Users table



Shipping\_Addresses table



Products table



Purchases table



# Keys

In a relational database, all rows must be unique. The column or combination of columns that make it unique are referred to as **key(s)**.

- **Natural Key:** From real world data, SSN's, customer account numbers, driver license numbers
- **Surrogate Key:** Keys artificially created by an application to make a row unique



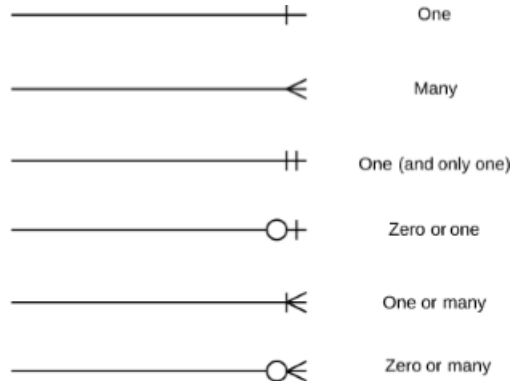
# Keys

In a relational database, all rows must be unique. The column or combination of columns that make it unique are referred to as **key(s)**.

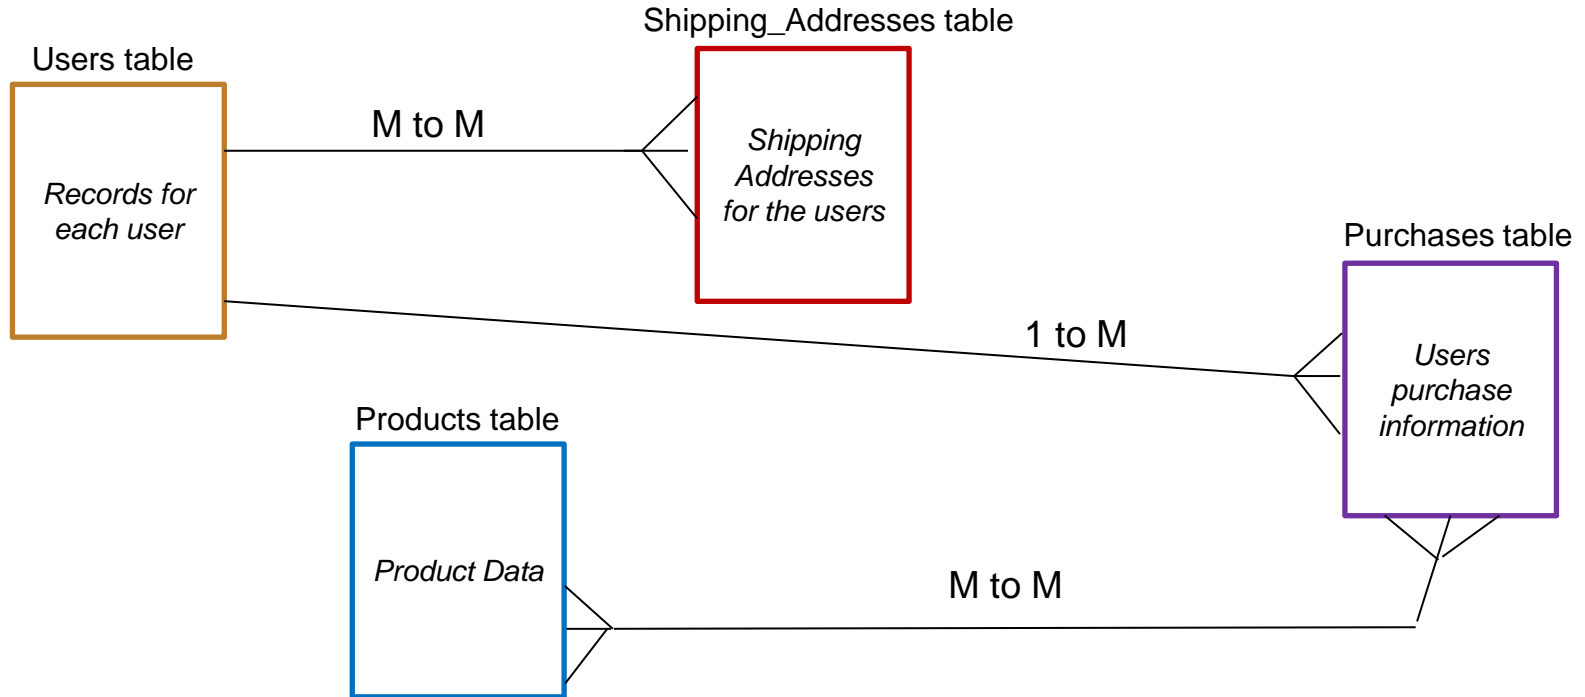
- **Primary Key:** column or columns in a table that uniquely identify the row. These cannot be duplicated.
  - If you say that SSN is your key, there cannot be more than one row with the same SSN.
- **Foreign Key:** A key that exists in another table, in which the latter is the primary key.

# Cardinality


- Describes relationship between two tables
- Relationship between a row in one table and a row of another table.
- Options are one or many
- 1 to 1, 1 to M, M to M




# Amazon Scenario




# Other examples



a one-to-one relationship




a one-to-many relationship



a many-to-many relationship

# Many to Many Relationship

Many records in one table  
relate to  
Many records in another table



# Student and Class Tables

student

ID	Name
1	John
2	Mark
3	Sarah
4	Claire

class

ID	Name
2	DB01
5	PH01
7	WEB01
8	WEB02

# Student and Class Tables

student

ID	Name	Class ID
1	John	2, 5
2	Mark	5, 7
3	Sarah	8
4	Claire	2, 5, 8

class


ID	Name	Student ID
2	DB01	1, 4
5	PH01	1, 2, 4
7	WEB01	2
8	WEB02	3, 4

# Joining Table

Instead of this:




We have this:






# The Joining Table



Captures every instance of student and class



student

ID	Name
1	John
2	Mark
3	Sarah
4	Claire

student\_class

Student ID	Class ID
1	2
1	5
2	5
2	7
3	8
4	2
4	5
4	8

class


ID	Name
2	DB01
5	PH01
7	WEB01
8	WEB02

# Joins


Joins in SQL allow us to pull in data from several tables.

A JOIN is an INNER JOIN


## SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

© C.L. Moffatt, 2008

# Joins : Inner Join or Join

An inner join returns the rows in Table A that has a matching key value in Table B, the Venn Diagram representation is as follows:



```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```

# Joins : Inner Join Example

Consider the following example: **We need a SQL query that returns all the cities in Texas. In my result, I want to see the full state name (not the code) followed by the cities in that state..**

- The city table contains the list of cities by state abbreviation... but it is missing the full state name.
- The state table contains the list of all states, but it has no data for cities.

city_id	city_name	state_abbreviation	population	area
1	Abilene	TX	123420	276.4
2	Akron	OH	197597	160.6
3	Albany	NY	96460	56.8
4	Albuquerque	NM	560513	487.4
5	Alexandria	VA	159428	38.8
6	Allen	TX	105623	70.2
7	Allentown	PA	121442	45.3
8	Amarillo	TX	199371	262.6
9	Anaheim	CA	350365	129.5
10	Anchorage	AK	288000	4420.1
11	Ann Arbor	MI	119980	72.8

state_abbreviation	state_name	population	area	capital	sales_tax	state_nickname	census_region
AL	Alabama	4903185	135767	198	4.000	Heart of Dixie	South
AK	Alaska	731545	1723337	155	0.000	The Last Frontier	West
AZ	Arizona	7278717	295234	238	5.600	Grand Canyon State	West
AR	Arkansas	3017804	137732	175	6.500	The Natural State	South
CA	California	39512223	423967	264	7.250	The Golden State	West
CO	Colorado	5736736	269601	88	2.900	Centennial State	West
CT	Connecticut	3565278	14357	133	6.350	Constitution State	Northeast
RI	Rhode Island	1097194	1545	01	0.000	The First State	South


- What we need to do is combine both tables:
  - Fortunately, these two tables are “related” via the state abbreviation value. Both tables refer to the column as state\_abbreviation.


# Joins : Inner Join Example


We can combine all of these facts to write a query that combines these two tables:

```
SELECT
FROM city c
INNER JOIN state s
ON c.state_abbreviation =
s.state_abbreviation
WHERE c.state_abbreviation =
'TX';
```

	Data Output	Explain	Messages	Notifications
	city_name character varying (50)		state_name character varying (50)	
1	Abilene		Texas	
2	Allen		Texas	
3	Amarillo		Texas	
4	Arlington		Texas	
5	Austin		Texas	
6	Beaumont		Texas	
7	Brownsville		Texas	
8	Carrollton		Texas	
9	College Station		Texas	
10	Corpus Christi		Texas	
11	Dallas		Texas	
12	Denton		Texas	
13	Edinburg		Texas	
14	El Paso		Texas	

luxury_cars	
 l_id	INTEGER
luxury_car_names	CHARACTER VARYING(250)

sports_cars	
 s_id	INTEGER
sports_car_names	CHARACTER VARYING(250)



```

24
25 SELECT L_ID, luxury_car_names, S_ID, sports_car_names
26 FROM Luxury_cars
27 INNER JOIN Sports_cars
28 ON luxury_car_names= sports_car_names;
29
30
31

```

```

24
25 SELECT L_ID, luxury_car_names, S_ID, sports_car_names
26 FROM Luxury_cars
27 JOIN Sports_cars
28 ON luxury_car_names= sports_car_names;
29
30
31

```

31:1 [712] INS

Log

1: Luxury\_cars [2] ×








* l_id	luxury_car_names	s_id	sports_car_names
1	1 Chevrolet Corvette	3	3 Chevrolet Corvette
2	2 Mercedes Benz SL Class	4	4 Mercedes Benz SL Class

<https://www.javatpoint.com/postgresql-join>




Let's write some inner join queries!



# Joins : Left Outer Join (can also be called Left Join)

The Left Outer Join returns all the rows on the “left” side table of the join, it will attempt to match to the right side. If there is match... If it can't find a match it includes it in the result, but with NULL values.

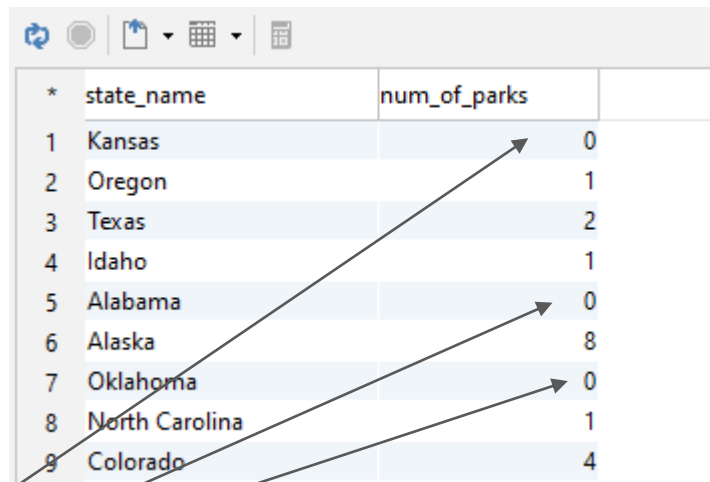


```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```

# Joins : Left Outer Join Example

```
SELECT state_name, COUNT(park_state.park_id) AS  
num_of_parks  
FROM state  
LEFT JOIN park_state  
ON state.state_abbreviation =  
park_state.state_abbreviation  
GROUP BY state_name;
```

Note that the state\_nicknames for Kansas, Alabama, and Oklahoma don't have any parks in the database, so the Left Outer Join instead creates these 0 placeholders.

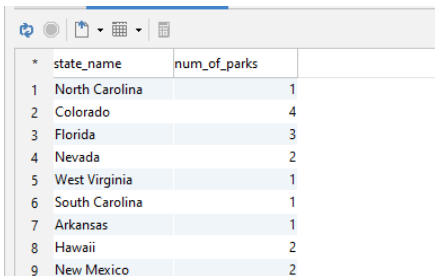


*	state_name	num_of_parks
1	Kansas	0
2	Oregon	1
3	Texas	2
4	Idaho	1
5	Alabama	0
6	Alaska	8
7	Oklahoma	0
8	North Carolina	1
9	Colorado	4

# Joins : Left Join vs Inner Join

With the same data set as the previous slide, let's compare the LEFT OUTER vs an INNER.

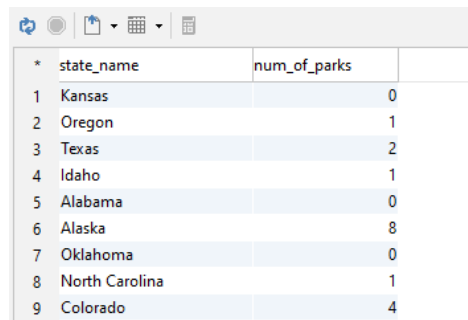
```
SELECT state_name, COUNT(park_state.park_id) AS  
num_of_parks  
FROM state  
JOIN park_state ON state.state_abbreviation =  
park_state.state_abbreviation  
GROUP BY state_name;
```




	state_name	num_of_parks
1	North Carolina	1
2	Colorado	4
3	Florida	3
4	Nevada	2
5	West Virginia	1
6	South Carolina	1
7	Arkansas	1
8	Hawaii	2
9	New Mexico	2


With the INNER JOIN, the rows for which there are no matches on the key are dropped from the final result set.


```
SELECT state_name, COUNT(park_state.park_id) AS  
num_of_parks  
FROM state  
LEFT JOIN park_state  
ON state.state_abbreviation =  
park_state.state_abbreviation  
GROUP BY state_name;
```



	state_name	num_of_parks
1	Kansas	0
2	Oregon	1
3	Texas	2
4	Idaho	1
5	Alabama	0
6	Alaska	8
7	Oklahoma	0
8	North Carolina	1
9	Colorado	4

luxury_cars	
 l_id	INTEGER
luxury_car_names	CHARACTER VARYING(250)

sports_cars	
 s_id	INTEGER
sports_car_names	CHARACTER VARYING(250)



```

31
32 SELECT L_ID, luxury_car_names, S_ID, sports_car_names
33 FROM Luxury_cars
34 LEFT OUTER JOIN Sports_cars
35 ON luxury_car_names= sports_car_names;
36
37

```

```

30
31
32 SELECT L_ID, luxury_car_names, S_ID, sports_car_names
33 FROM Luxury_cars
34 LEFT JOIN Sports_cars
35 ON luxury_car_names= sports_car_names;
36
37


```


15:36 [853] INS

Log 1: Luxury\_cars [5] ×


* l_id	luxury_car_names	s_id	sports_car_names
1	1 Chevrolet Corvette	3	Chevrolet Corvette
2	2 Mercedes Benz SL Class	4	Mercedes Benz SL Class
3	3 Audi A7	(null)	(null)
4	4 Genesis G90	(null)	(null)
5	5 Lincoln Continental	(null)	(null)

<https://www.javatpoint.com/postgresql-join>

luxury_cars	
 l_id	INTEGER
luxury_car_names	CHARACTER VARYING(250)

sports_cars	
 s_id	INTEGER
sports_car_names	CHARACTER VARYING(250)

**Left Outer Join: Select  
Rows From the Left table**



```

45
46 SELECT L_ID, luxury_car_names, S_ID, sports_car_names
47 FROM Luxury_cars
48 LEFT JOIN Sports_cars
49 ON luxury_car_names= sports_car_names
50 WHERE S_ID IS NULL;

```

50:22 [1191] INS

Log 1: Luxury\_cars [3] ×

* l_id	luxury_car_names	s_id	sports_car_names
1	3 Audi A7	(null)	(null)
2	4 Genesis G90	(null)	(null)
3	5 Lincoln Continental	(null)	(null)

<https://www.javatpoint.com/postgresql-join>

# Unions

A union is a combination of two result sets. The following pattern is used:

[SQL Query 1]

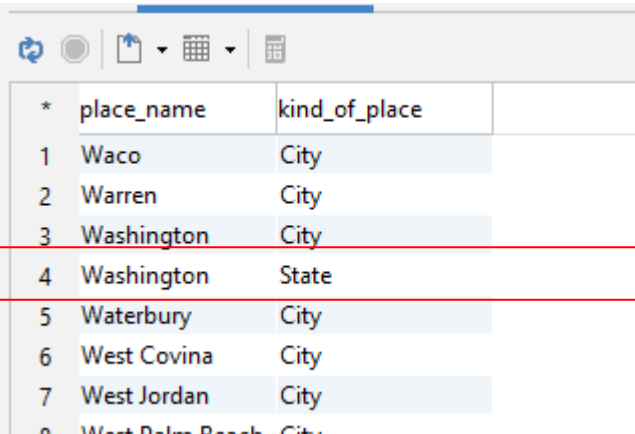
**UNION**

[SQL Query 2]

# Unions Example

Consider the following query:

```
SELECT city_name AS place_name FROM city WHERE city_name LIKE 'W%'  
UNION  
SELECT state_name AS place_name FROM state WHERE state_name LIKE 'W%'  
ORDER BY place_name;
```



The screenshot shows a database query result with two columns: 'place\_name' and 'kind\_of\_place'. The results are ordered alphabetically by 'place\_name'. The first three rows are cities: Waco, Warren, and Washington. The fourth row is the state of Washington, which is highlighted with a red rectangle. The following rows are Waterbury, West Covina, West Jordan, and West Palm Beach.

	place_name	kind_of_place
*	place_name	kind_of_place
1	Waco	City
2	Warren	City
3	Washington	City
4	Washington	State
5	Waterbury	City
6	West Covina	City
7	West Jordan	City
8	West Palm Beach	City

Result of query first query  
(returns the City)

Result of query second query  
(returns the State)



# Union All

## 2) PostgreSQL UNION ALL example

The following statement uses the `UNION ALL` operator to combine result sets from the `topRatedFilms` and `mostPopularFilms` tables:

```
SELECT * FROM topRatedFilms
UNION ALL
SELECT * FROM mostPopularFilms;
```

	title character varying	release_year smallint
1	The Shawshank Redemption	1994
2	The Godfather	1972
3	12 Angry Men	1957
4	An American Pickle	2020
5	The Godfather	1972
6	Greyhound	2020

In this example, the duplicate row is retained in the result set.